

Java Platform, Standard Edition

Internationalization Guide



Release 11
E94896-02
November 2018

ORACLE®

Copyright © 1993, 2018, Oracle and/or its affiliates. All rights reserved.

This software and related documentation are provided under a license agreement containing restrictions on use and disclosure and are protected by intellectual property laws. Except as expressly permitted in your license agreement or allowed by law, you may not use, copy, reproduce, translate, broadcast, modify, license, transmit, distribute, exhibit, perform, publish, or display any part, in any form, or by any means. Reverse engineering, disassembly, or decompilation of this software, unless required by law for interoperability, is prohibited.

The information contained herein is subject to change without notice and is not warranted to be error-free. If you find any errors, please report them to us in writing.

If this is software or related documentation that is delivered to the U.S. Government or anyone licensing it on behalf of the U.S. Government, then the following notice is applicable:

U.S. GOVERNMENT END USERS: Oracle programs, including any operating system, integrated software, any programs installed on the hardware, and/or documentation, delivered to U.S. Government end users are "commercial computer software" pursuant to the applicable Federal Acquisition Regulation and agency-specific supplemental regulations. As such, use, duplication, disclosure, modification, and adaptation of the programs, including any operating system, integrated software, any programs installed on the hardware, and/or documentation, shall be subject to license terms and license restrictions applicable to the programs. No other rights are granted to the U.S. Government.

This software or hardware is developed for general use in a variety of information management applications. It is not developed or intended for use in any inherently dangerous applications, including applications that may create a risk of personal injury. If you use this software or hardware in dangerous applications, then you shall be responsible to take all appropriate fail-safe, backup, redundancy, and other measures to ensure its safe use. Oracle Corporation and its affiliates disclaim any liability for any damages caused by use of this software or hardware in dangerous applications.

Oracle and Java are registered trademarks of Oracle and/or its affiliates. Other names may be trademarks of their respective owners.

Intel and Intel Xeon are trademarks or registered trademarks of Intel Corporation. All SPARC trademarks are used under license and are trademarks or registered trademarks of SPARC International, Inc. AMD, Opteron, the AMD logo, and the AMD Opteron logo are trademarks or registered trademarks of Advanced Micro Devices. UNIX is a registered trademark of The Open Group.

This software or hardware and documentation may provide access to or information about content, products, and services from third parties. Oracle Corporation and its affiliates are not responsible for and expressly disclaim all warranties of any kind with respect to third-party content, products, and services unless otherwise set forth in an applicable agreement between you and Oracle. Oracle Corporation and its affiliates will not be responsible for any loss, costs, or damages incurred due to your access to or use of third-party content, products, or services, except as set forth in an applicable agreement between you and Oracle.

Contents

Preface

Audience	vi
Documentation Accessibility	vi
Related Documents	vi
Conventions	vi

1 Internationalization Enhancements

Internationalization Enhancements in JDK 11	1-1
Unicode 10.0.0	1-1
Internationalization Enhancements in JDK 10	1-1
Additional Unicode Language-Tag Extensions	1-1
Internationalization Enhancements in JDK 9	1-2
Unicode 8.0	1-2
CLDR Locale Data Enabled by Default	1-2
UTF-8 Properties Files	1-3

2 Internationalization Overview

Text Representation	2-1
Locale Identification and Localization	2-2
Locales	2-2
Locale Class	2-2
Supported Locales	2-2
Localized Resources	2-3
ResourceBundle Class	2-3
ListResourceBundle Class	2-3
PropertyResourceBundle Class	2-3
Date and Time Handling	2-4
Text Processing	2-4
Formatting	2-4
Format Class	2-4
DateFormat Class	2-4

SimpleDateFormat Class	2-5
DateFormatSymbols Class	2-5
NumberFormat Class	2-5
DecimalFormat Class	2-5
DecimalFormatSymbols Class	2-6
ChoiceFormat Class	2-6
MessageFormat Class	2-6
ParsePosition Class	2-6
FieldPosition Class	2-6
Locale-Sensitive String Operations	2-6
Collator Class	2-7
RuleBasedCollator Class	2-7
CollationElementIterator Class	2-7
CollationKey Class	2-7
BreakIterator Class	2-7
StringCharacterIterator Class	2-8
CharacterIterator Interface	2-8
Normalizer Class	2-8
Locale-Sensitive Services SPIs	2-8
Character Encoding Conversion	2-8
Supported Encodings	2-9
Stream I/O	2-9
Reader and Writer Classes	2-9
PrintStream Class	2-9
Charset Package	2-9
Input Methods	2-9
Input Method Support in Swing	2-10
Input Method Framework	2-10

3 Supported Encodings

Basic Encoding Set (contained in java.base module)	3-1
Extended Encoding Set (contained in jdk.charsets module)	3-3

4 Supported Calendars

5 Supported Fonts

Support for Physical Fonts	5-1
Support for Logical Fonts	5-1

6 Font Configuration Files

Supported Platforms	6-1
Loading Font Configuration Files	6-1
Names Used in Font Configuration Files	6-2
Properties for All Platforms	6-3
Version Property	6-3
Component Font Mappings	6-3
Search Sequences	6-3
Exclusion Ranges	6-5
Proportional Fonts	6-5
Font File Names	6-6
Appended Font Path	6-6
Properties for Windows	6-6
Property for Solaris and Linux	6-7

Preface

This guide summarizes the internationalization APIs and features of the Java SE Platform.

Audience

This guide is intended for Java programmers who want to design applications so that they can be adapted to various languages and regions without engineering changes.

Documentation Accessibility

For information about Oracle's commitment to accessibility, visit the Oracle Accessibility Program website at <http://www.oracle.com/pls/topic/lookup?ctx=acc&id=docacc>.

Access to Oracle Support

Oracle customers that have purchased support have access to electronic support through My Oracle Support. For information, visit <http://www.oracle.com/pls/topic/lookup?ctx=acc&id=info> or visit <http://www.oracle.com/pls/topic/lookup?ctx=acc&id=trs> if you are hearing impaired.

Related Documents

For coding examples and step-by-step instructions, see the [Internationalization Trail](#) in The Java Tutorials (Java SE 8 and earlier).

Conventions

The following text conventions are used in this document:

Convention	Meaning
boldface	Boldface type indicates graphical user interface elements associated with an action, or terms defined in text or the glossary.
<i>italic</i>	Italic type indicates book titles, emphasis, or placeholder variables for which you supply particular values.
monospace	Monospace type indicates commands within a paragraph, URLs, code in examples, text that appears on the screen, or text that you enter.

1

Internationalization Enhancements

Recent releases of the JDK include enhancements to the internationalization process to support updated standards.

Topics:

- [Internationalization Enhancements in JDK 11](#)
- [Internationalization Enhancements in JDK 10](#)
- [Internationalization Enhancements in JDK 9](#)

Internationalization Enhancements in JDK 11

Internationalization enhancements for JDK 11 include:

- [Unicode 10.0.0](#)

Unicode 10.0.0

Support has been added for Unicode 10.0.0. Java Platform, Standard Edition (Java SE) 9 and 10 supported Unicode 8.0.

The Unicode [10.0](#) standard includes 16,018 characters and 10 scripts that were introduced since Unicode 8.0, all of which are supported in Java SE 11.

Internationalization Enhancements in JDK 10

Internationalization enhancements for JDK 10 include:

- [Additional Unicode Language-Tag Extensions](#)

Additional Unicode Language-Tag Extensions

The IETF BCP (best current practice) 47 language tags standard, which has been supported in the `Locale` class since Java SE 7, includes a Unicode extension subtag. As of Java SE 9, only the `-ca` (calendar) and `-nu` (number) extensions are supported.

Java SE 10 adds support for the following additional extensions in the relevant JDK classes:

- `-cu` (currency type)
- `-fw` (first day of week)
- `-rg` (region override)
- `-tz` (time zone)

In JDK 10, if an application specifies a locale of `en-US-u-cu-EUR`, which means US English with Euro currency, `java.util.Currency.getInstance(locale)` instantiates a Euro `Currency`. If the locale is `en-US-u-cu-JPY`, a Japanese Yen `Currency` is instantiated.

Internationalization Enhancements in JDK 9

Internationalization enhancements for Oracle Java Development Kit 9 include:

- [Unicode 8.0](#)
- [CLDR Locale Data Enabled by Default](#)
- [UTF-8 Properties Files](#)

Unicode 8.0

Support has been added for Unicode 8.0. Java Platform, Standard Edition (Java SE) 8 supported Unicode 6.2.

The Unicode [6.3](#), [7.0](#), and [8.0](#) standards introduced 10,555 characters, 29 scripts, and 42 blocks, all of which are supported in Java SE 9.

CLDR Locale Data Enabled by Default

The XML-based locale data of the Unicode Common Locale Data Repository (CLDR), first added in JDK 8, is the default locale data in JDK 9. In previous releases, the default was `JRE`.

There are four distinct sources for locale data, identified by the following keywords:

- `CLDR` represents the locale data provided by the Unicode [CLDR](#) project.
- `HOST` represents the current user's customization of the underlying operating system's settings. It works only with the user's default locale, and the customizable settings may vary depending on the operating system. However, primarily date, time, number, and currency formats are supported.
- `SPI` represents the locale-sensitive services implemented by the installed Service Provider Interface (SPI) providers.
- `COMPAT` (formerly called `JRE`) represents the locale data that is compatible with releases prior to JDK 9. `JRE` can still be used as the value, but `COMPAT` is preferred.

To select a locale data source, use the `java.locale.providers` system property, listing the data sources in the preferred order. If a provider cannot offer the requested locale data, the search proceeds to the next provider in order. For example:

```
java.locale.providers=HOST,SPI,CLDR,COMPAT
```

If you do not set this property, the default behavior is equivalent to the following setting:

```
java.locale.providers=CLDR,COMPAT,SPI
```

To enable behavior that is compatible with JDK 8, set the `java.locale.providers` system property to a value with `COMPAT` to the left of `CLDR`.

See the [JDK 9 and JRE 9 Supported Locales](#) page for supported locales. See [java.util.spi.LocaleServiceProvider API specification](#) for the related API.

UTF-8 Properties Files

In Java SE 9, properties files are loaded in UTF-8 encoding. In previous releases, ISO-8859-1 encoding was used for loading property resource bundles. UTF-8 is a much more convenient way to represent non-Latin characters.

Most existing properties files should not be affected: UTF-8 and ISO-8859-1 have the same encoding for ASCII characters, and human-readable non-ASCII ISO-8859-1 encoding is not valid UTF-8. If an invalid UTF-8 byte sequence is detected, the Java runtime automatically rereads the file in ISO-8859-1.

If there is an issue, consider the following options:

- Convert the properties file into UTF-8 encoding.
- Specify the runtime system property for the properties file's encoding, as in this example:

```
java.util.PropertyResourceBundle.encoding=ISO-8859-1
```

See [java.util.PropertyResourceBundle](#).

2

Internationalization Overview

Internationalization is the process of designing an application so that it can be adapted to various languages and regions without engineering changes. Sometimes the term internationalization is abbreviated as i18n, because there are 18 letters between the first "i" and the last "n."

An internationalized program has the following characteristics:

- With the addition of localization data, the same executable can run worldwide.
- Textual elements, such as status messages and the GUI component labels, are not hardcoded in the program. Instead they are stored outside the source code and retrieved dynamically.
- Support for new languages does not require recompilation.
- Culturally-dependent data, such as dates and currencies, appear in formats that conform to the end user's region and language.
- It can be localized quickly.

The global Internet demands global software - that is, software that can be developed independently of the countries or languages of its users, and then localized for multiple countries or regions. The Java Platform provides a rich set of APIs for developing global applications. These internationalization APIs are based on the Unicode standard and include the ability to adapt text, numbers, dates, currency, and user-defined objects to any country's conventions.

This guide summarizes the internationalization APIs and features of the Java Platform, Standard Edition. For coding examples and step-by-step instructions, see the [Internationalization Trail](#) in the Java Tutorials.

Text Representation

The Java programming language is based on the [Unicode](#) character set, and several libraries implement the Unicode standard. Unicode is an international character set standard which supports all of the major scripts of the world, as well as common technical symbols. The original Unicode specification defined characters as fixed-width 16-bit entities, but the Unicode standard has since been changed to allow for characters whose representation requires more than 16 bits. The range of legal code points is now U+0000 to U+10FFFF. An encoding defined by the standard, UTF-16, allows to represent all Unicode code points using one or two 16-bit units.

The primitive data type `char` in the Java programming language is an unsigned 16-bit integer that can represent a Unicode code point in the range U+0000 to U+FFFF, or the code units of UTF-16. The various types and classes in the Java platform that represent character sequences - `char[]`, implementations of `java.lang.CharSequence` (such as the `String` class), and implementations of `java.text.CharacterIterator` - are UTF-16 sequences. Most Java source code is written in ASCII, a 7-bit character encoding, or ISO-8859-1, an 8-bit character encoding, but is translated into UTF-16 before processing.

The `Character` class is an object wrapper for the `char` primitive type. The `Character` class also contains static methods such as `isLowerCase()` and `isDigit()` for determining the properties of a character. These methods have overloads that accept either a `char` (which allows representation of Unicode code points in the range U+0000 to U+FFFF) or an `int` (which allows representation of all Unicode code points).

Locale Identification and Localization

A `Locale` object is an identifier for a particular combination of language and region. Localization is the process of adapting software for a specific region or language by adding locale-specific components and translating text.

Locales

On the Java platform, a locale is simply an identifier for a particular combination of language and region. It is not a collection of locale-specific attributes. Instead, each locale-sensitive class maintains its own locale-specific information. With this design, there is no difference in how user and system objects maintain their locale-specific resources. Both use the standard localization mechanism.

Java programs are *not* assigned a single global locale. All locale-sensitive operations may be explicitly given a locale as an argument. This greatly simplifies multilingual programs. While a global locale is not enforced, a default locale is available for programs that do not wish to manage locales explicitly. A default locale also makes it possible to affect the behavior of the entire presentation with a single choice.

Java locales act as requests for certain behavior from another object. For example, a French Canadian locale passed to a `Calendar` object asks that the `Calendar` behave correctly for the customs of Quebec. It is up to the object accepting the locale to do the right thing. If the object has not been localized for a particular locale, it will try to find a "close" match with a locale for which it has been localized. Thus if a `Calendar` object was not localized for French Canada, but was localized for the French language in general, it would use the French localization instead.

Locale Class

A `Locale` object represents a specific geographical, political, or cultural region. An operation that requires a locale to perform its task is called locale-sensitive and uses the `Locale` object to tailor information for the user. For example, displaying a number is a locale-sensitive operation - the number should be formatted according to the customs and conventions of the user's native country, region, or culture.

Supported Locales

On the Java Platform, there does not have to be a single set of supported locales, since each class maintains its own localizations. Nevertheless, there is a consistent set of localizations supported by the classes of the Java Platform. Other implementations of the Java Platform may support different locales. Locales that are supported by the JDK and JRE are summarized by release, using the search field for the [Oracle Technology Network](#) page, search for "Supported Locales".

Localized Resources

All locale-sensitive classes must be able to access resources customized for the locales they support. To aid in the process of localization, it helps to have these resources grouped together by locale and separated from the locale-neutral parts of the program.

ResourceBundle Class

The class `ResourceBundle` is an abstract base class representing containers of resources. Programmers create subclasses of `ResourceBundle` that contain resources for a particular locale. New resources can be added to an instance of `ResourceBundle`, or new instances of `ResourceBundle` can be added to a system without affecting the code that uses them. Packaging resources as classes allows developers to take advantage of Java's class loading mechanism to find resources.

Resource bundles contain locale-specific objects. When a program needs a locale-specific resource, such as a `String` object, the program can load it from the resource bundle that is appropriate for the current user's locale. In this way, the programmer can write code that is largely independent of the user's locale, isolating most, if not all, of the locale-specific information in resource bundles.

This allows Java programmers to write code that can:

- be easily localized, or translated, into different languages
- handle multiple locales at once
- be easily modified later to support even more locales

ResourceBundle.Control Class

`ResourceBundle.Control` is a nested class of `ResourceBundle`. It defines methods to be called by the `ResourceBundle.getBundle` factory methods so that the resource bundle loading behavior may be changed. For example, application specific resource bundle formats, such as XML, could be supported by overriding the methods.

Since Java SE 9, `ResourceBundle.Control` is not supported in named modules. Existing code using `Control` is expected to work, but for new code in a named module, implement `basenameProvider` and load the resource bundle from there. See [Resource Bundles and Named Modules](#).

ListResourceBundle Class

`ListResourceBundle` is an abstract subclass of `ResourceBundle` that manages resources for a locale in a convenient and easy to use list.

PropertyResourceBundle Class

`PropertyResourceBundle` is a concrete subclass of `ResourceBundle` that manages resources for a locale using a set of static strings from a property file.

Date and Time Handling

The Date-Time package, `java.time`, introduced in Java SE 8, provides a comprehensive model for date and time. Although `java.time` is based on the International Organization for Standardization (ISO) calendar system, commonly used global calendars are also supported.

See [The Date-Time Packages](#) lesson in The Java Tutorials (Java SE 8 and earlier).

Text Processing

Text processing involves formatting locale-sensitive information such as, currencies, dates, times, and text messages. It also includes manipulating text in a locale-sensitive manner, meaning that string operations, such as searching and sorting, are properly performed regardless of locale.

Formatting

It is in formatting data for output that many cultural conventions are applied. Numbers, dates, times, and messages may all require formatting before they can be displayed. The Java platform provides a set of flexible formatting classes that can handle both the standard locale formats and programmer defined custom formats. These formatting classes are also able to parse formatted strings back into their constituent objects.

Format Class

The class `Format` is an abstract base class for formatting locale-sensitive information such as dates, times, messages, and numbers. Three main subclasses are provided: `DateFormat`, `NumberFormat`, and `MessageFormat`. These three also provide subclasses of their own.

DateFormat Class

Dates and times are stored internally in a locale-independent way, but should be formatted so that they can be displayed in a locale-sensitive manner. For example, the same date might be formatted as:

- November 3, 1997 (English)
- 3 novembre 1997 (French)

The class `DateFormat` is an abstract base class for formatting and parsing date and time values in a locale-independent manner. It has a number of static factory methods for getting standard time formats for a given locale.

The `DateFormat` object uses `Calendar` and `TimeZone` objects in order to interpret time values. By default, a `DateFormat` object for a given locale will use the appropriate `Calendar` object for that locale and the system's default `TimeZone` object. The programmer can override these choices if desired.

SimpleDateFormat Class

The class `SimpleDateFormat` is a concrete class for formatting and parsing dates and times in a locale-sensitive manner. It allows for formatting (milliseconds to text), parsing (text to milliseconds), and normalization.

DateFormatSymbols Class

The class `DateFormatSymbols` is used to encapsulate localizable date-time formatting data, such as the names of the months, the names of the days of the week, time of day, and the time zone data. The `DateFormat` and `SimpleDateFormat` classes both use the `DateFormatSymbols` class to encapsulate this information.

Usually, programmers will not use the `DateFormatSymbols` directly. Rather, they will implement formatting with the `DateFormat` class's factory methods.

NumberFormat Class

The class `NumberFormat` is an abstract base class for formatting and parsing numeric data. It contains a number of static factory methods for getting different kinds of locale-specific number formats.

The `NumberFormat` class helps programmers to format and parse numbers for any locale. Code using this class can be completely independent of the locale conventions for decimal points, thousands-separators, the particular decimal digits used, or whether the number format is even decimal. The application can also display a number as a normal decimal number, currency, or percentage:

- 1,234.5 (decimal number in U.S. format)
- \$1,234.50 (U.S. currency in U.S. format)
- 1.234,50 € (European currency in German format)
- 123.450% (percent in German format)

DecimalFormat Class

Numbers are stored internally in a locale-independent way, but should be formatted so that they can be displayed in a locale-sensitive manner. For example, when using `"#,###.00"` as a pattern, the same number might be formatted as:

- 1.234,56 (German)
- 1,234.56 (English)

The class `DecimalFormat`, which is a concrete subclass of the `NumberFormat` class, can format decimal numbers. Programmers generally will not instantiate this class directly but will use the factory methods provided.

The `DecimalFormat` class has the ability to take a pattern string to specify how a number should be formatted. The pattern specifies attributes such as the precision of the number, whether leading zeros should be printed, and what currency symbols are used. The pattern string can be altered if a program needs to create a custom format.

DecimalFormatSymbols Class

The class `DecimalFormatSymbols` represents the set of symbols (such as the decimal separator, the grouping separator, and so on) needed by `DecimalFormat` to format numbers. `DecimalFormat` creates for itself an instance of `DecimalFormatSymbols` from its locale data. A programmer needing to change any of these symbols can get the `DecimalFormatSymbols` object from the `DecimalFormat` object and then modify it.

ChoiceFormat Class

The class `ChoiceFormat` is a concrete subclass of the `NumberFormat` class. The `ChoiceFormat` class allows the programmer to attach a format to a range of numbers. It is generally used in a `MessageFormat` object for handling plurals.

MessageFormat Class

Programs often need to build messages from sequences of strings, numbers and other data. For example, the text of a message displaying the number of files on a disk drive will vary:

- The disk C contains 100 files.
- The disk D contains 1 file.
- The disk F contains 0 files.

If a message built from sequences of strings and numbers is hard-coded, it cannot be translated into other languages. For example, note the different positions of the parameters "3" and "G" in the following translations:

- The disk G contains 3 files. (English)
- Il y a 3 fichiers sur le disque G. (French)

The class `MessageFormat` provides a means to produce concatenated messages in language-neutral way. The `MessageFormat` object takes a set of objects, formats them, and then inserts the formatted strings into the pattern at the appropriate places.

ParsePosition Class

The class `ParsePosition` is used by the `Format` class and its subclasses to keep track of the current position during parsing. The `parseObject()` method in the `Format` class requires a `ParsePosition` object as an argument.

FieldPosition Class

The `FieldPosition` class is used by the `Format` class and its subclasses to identify fields in formatted output. One version of the `format()` method in the `Format` class requires a `FieldPosition` object as an argument.

Locale-Sensitive String Operations

Programs frequently need to manipulate strings. Common operations on strings include searching and sorting. Some tasks, such as collating strings or finding various

boundaries in text, are surprisingly difficult to get right and are even more difficult when multiple languages must be considered. The Java Platform provides classes for handling many of these common string manipulations in a locale-sensitive manner.

Collator Class

The `Collator` class performs locale-sensitive string comparison. Programmers use this class to build searching and alphabetical sorting routines for natural language text. `Collator` is an abstract base class. Its subclasses implement specific collation strategies. One subclass, `RuleBasedCollator`, is applicable to a wide set of languages. Other subclasses may be created to handle more specialized needs.

RuleBasedCollator Class

The `RuleBasedCollator` class, which is a concrete subclass of the `Collator` class, provides a simple, data-driven, table collator. Using `RuleBasedCollator`, a programmer can create a customized table-based collator. For example, a programmer can build a collator that will ignore (or notice) uppercase letters, accents, and Unicode combining characters.

CollationElementIterator Class

The `CollationElementIterator` class is used as an iterator to walk through each character of an international string. Programmers use the iterator to return the ordering priority of the positioned character. The ordering priority of a character, or key, defines how a character is collated in the given `Collator` object. The `CollationElementIterator` class is used by the `compare()` method of the `RuleBasedCollator` class.

CollationKey Class

A `CollationKey` object represents a string under the rules of a specific `Collator` object. Comparing two `CollationKey` objects returns the relative order of the strings they represent. Using `CollationKey` objects to compare strings is generally faster than using the `Collator.compare()` method. Thus, when the strings must be compared multiple times, for example when sorting a list of strings, it is more efficient to use `CollationKey` objects.

BreakIterator Class

The `BreakIterator` class indirectly implements methods for finding the position of the following types of boundaries in a string of text:

- potential line break
- sentence
- word
- character

The conventions on where to break lines, sentences, words, and characters vary from one language to another. Since the `BreakIterator` class is locale-sensitive, it can be used by programs that perform text operations. For example, consider a word processing program that can highlight a character, cut a word, move the cursor to the next sentence, or word-wrap at a line ending. This word processing program would

use break iterators to determine the logical boundaries in text, enabling it to perform text operations in a locale-sensitive manner.

StringCharacterIterator Class

The `StringCharacterIterator` class provides the ability to iterate over a string of Unicode characters in a bidirectional manner. This class uses a cursor to move within a range of text, and can return individual characters or their index values. The `StringCharacterIterator` class implements the character iterator functionality of the `CharacterIterator` interface.

CharacterIterator Interface

The `CharacterIterator` interface defines a protocol for bidirectional iteration over Unicode characters. Classes should implement this interface if they want to move about within a range of text and return individual Unicode characters or their index values. `CharacterIterator` is useful when performing character searches.

Normalizer Class

The `Normalizer` class provides methods to transform Unicode text into an equivalent composed or decomposed form. The class supports the *Unicode Normalization Forms* defined by the Unicode standard.

Locale-Sensitive Services SPIs

Locale sensitive services provided by classes in the `java.text` and `java.util` packages can be extended by implementing locale-sensitive services SPIs for locales the Java runtime has not yet supported.

Although JDK 9 no longer supports the extension mechanism, SPI implementations for internationalization functions in the `java.text.spi`, `java.util.spi`, and `java.awt.im.spi` packages will be loaded from the application's classpath.

In addition to localized symbols or names for the `Currency`, `Locale`, and `TimeZone` classes in the `java.util` package, implementations of the following classes in the `java.text` package can be plugged in with the SPIs.

- `BreakIterator`
- `Collator`
- `DateFormat`
- `DateFormatSymbols`
- `DecimalFormatSymbols`
- `NumberFormat`

Character Encoding Conversion

The Java platform uses Unicode as its native character encoding; however, many Java programs still need to handle text data in other encodings. Java therefore provides a set of classes that convert many standard character encodings to and from

Unicode. Java programs that need to deal with non-Unicode text data convert that data into Unicode, process the data as Unicode, then convert the result back to the external character encoding. The `InputStreamReader` and `OutputStreamWriter` classes provide methods that can convert between other character encodings and Unicode.

Supported Encodings

The `InputStreamReader`, `OutputStreamWriter`, and `String` classes can convert between Unicode and the set of character encodings listed in [Supported Encodings](#).

Stream I/O

The Java Platform provides features in the `java.io` package to improve the handling of character data: the `Reader` and `Writer` classes, and an enhancement to the `PrintStream` class.

Reader and Writer Classes

The `Reader` and `Writer` class hierarchies provide the ability to perform I/O operations on character streams. These hierarchies parallel the `InputStream` and `OutputStream` class hierarchies, but operate on streams of characters rather than streams of bytes. Character streams make it easy to write programs that are not dependent upon a specific character encoding, and are therefore easier to internationalize. The `Reader` and `Writer` classes also have the ability to convert between Unicode and other character encodings.

PrintStream Class

The `PrintStream` class produces output using the system's default character encoding and line terminator. This change allows methods such as `System.out.println()` to act more reasonably with non-ASCII data.

Charset Package

The `java.nio.charset` package provides the underpinnings for character encoding conversion. Applications can use its classes to fine-tune the behavior of built-in character converters. Developers can also create custom converters for character encodings that are not supported by built-in character converters, using the `java.nio.charset.spi` package.

Input Methods

Input methods are software components that let the user enter text in ways other than simple typing on a keyboard. They are commonly used to enter Japanese, Chinese, or Korean - languages using thousands of different characters - on keyboards with far fewer keys. However, the Java platform also supports input methods for other languages and the use of entirely different input mechanisms, such as handwriting or speech recognition.

The Java platform enables the use of native input methods provided by the host operating system, such as Windows or Solaris, as well as the implementation and use of input methods written in the Java programming language.

The term input methods does not refer to class methods of the Java programming language.

Input Method Support in Swing

The Swing text components provide an integrated user interface for text input using input methods. Depending on the locale, one of two input styles is used. With on-the-spot (inline) input, the style used for most locales, the input methods insert the text directly into the text component while the text is being composed. With below-the-spot input, the style used for Chinese locales, a separate composition window is used, which is positioned automatically to be near the point where the text is to be inserted after being committed.

An application using Swing text components does not have to coordinate the interaction between the text components and input methods. However, it should call `InputContext.endComposition` when all text must be committed, such as when a document is saved or printed.

Input Method Framework

The input method framework enables the collaboration between text editing components and input methods in entering text. Programmers who develop text editing components or input methods use this framework. Other application developers generally make only minimal use of it. For example, they should call `InputContext.endComposition` when all text must be committed, such as when a document is saved or printed.

3

Supported Encodings

The `java.io.InputStreamReader`, `java.io.OutputStreamWriter`, `java.lang.String` classes, and classes in the `java.nio.charset` package can convert between Unicode and a number of other character encodings. The supported encodings vary between different implementations of the Java SE Platform. The class description for `java.nio.charset.Charset` lists the encodings that any implementation of the Java SE Platform is required to support.

The following tables show the encoding sets supported by this version of the Oracle Java SE Platform. The canonical names used by the `java.nio` APIs are in many cases not the same as those used in the `java.io` and `java.lang` APIs.

Basic Encoding Set (contained in `java.base` module)

Canonical Name for <code>java.nio</code> API	Canonical Name for <code>java.io</code> API and <code>java.lang</code> API	Description
IBM00858	Cp858	Variant of Cp850 with Euro character
IBM437	Cp437	MS-DOS United States, Australia, New Zealand, South Africa
IBM775	Cp775	PC Baltic
IBM850	Cp850	MS-DOS Latin-1
IBM852	Cp852	MS-DOS Latin-2
IBM855	Cp855	IBM Cyrillic
IBM857	Cp857	IBM Turkish
IBM862	Cp862	PC Hebrew
IBM866	Cp866	MS-DOS Russian
ISO-8859-1	ISO8859_1	ISO-8859-1, Latin Alphabet No. 1
ISO-8859-2	ISO8859_2	Latin Alphabet No. 2
ISO-8859-4	ISO8859_4	Latin Alphabet No. 4
ISO-8859-5	ISO8859_5	Latin/Cyrillic Alphabet
ISO-8859-7	ISO8859_7	Latin/Greek Alphabet (ISO-8859-7:2003)
ISO-8859-9	ISO8859_9	Latin Alphabet No. 5
ISO-8859-13	ISO8859_13	Latin Alphabet No. 7
ISO-8859-15	ISO8859_15	Latin Alphabet No. 9
KOI8-R	KOI8_R	KOI8-R, Russian
KOI8-U	KOI8_U	KOI8-U, Ukrainian

Canonical Name for java.nio API	Canonical Name for java.io API and java.lang API	Description
US-ASCII	ASCII	American Standard Code for Information Interchange
UTF-8	UTF8	Eight-bit Unicode (or UCS) Transformation Format
UTF-16	UTF-16	Sixteen-bit Unicode (or UCS) Transformation Format, byte order identified by an optional byte-order mark
UTF-16BE	UnicodeBigUnmarked	Sixteen-bit Unicode (or UCS) Transformation Format, big-endian byte order
UTF-16LE	UnicodeLittleUnmarked	Sixteen-bit Unicode (or UCS) Transformation Format, little-endian byte order
UTF-32	UTF_32	32-bit Unicode (or UCS) Transformation Format, byte order identified by an optional byte-order mark
UTF-32BE	UTF_32BE	32-bit Unicode (or UCS) Transformation Format, big-endian byte order
UTF-32LE	UTF_32LE	32-bit Unicode (or UCS) Transformation Format, little-endian byte order
x-UTF-32BE-BOM	UTF_32BE_BOM	32-bit Unicode (or UCS) Transformation Format, big-endian byte order, with byte-order mark
x-UTF-32LE-BOM	UTF_32LE_BOM	32-bit Unicode (or UCS) Transformation Format, little-endian byte order, with byte-order mark
windows-1250	Cp1250	Windows Eastern European
windows-1251	Cp1251	Windows Cyrillic
windows-1252	Cp1252	Windows Latin-1
windows-1253	Cp1253	Windows Greek
windows-1254	Cp1254	Windows Turkish
windows-1257	Cp1257	Windows Baltic
Not available	UnicodeBig	Sixteen-bit Unicode (or UCS) Transformation Format, big-endian byte order, with byte-order mark
x-IBM737	Cp737	PC Greek
x-IBM874	Cp874	IBM Thai
x-UTF-16LE-BOM	UnicodeLittle	Sixteen-bit Unicode (or UCS) Transformation Format, little-endian byte order, with byte-order mark

Extended Encoding Set (contained in `jdk.charsets` module)

Canonical Name for <code>java.nio</code> API	Canonical Name for <code>java.io</code> API and <code>java.lang</code> API	Description
Big5	Big5	Big5, Traditional Chinese
Big5-HKSCS	Big5_HKSCS	Big5 with Hong Kong extensions, Traditional Chinese (incorporating 2001 revision)
EUC-JP	EUC_JP	JISX 0201, 0208 and 0212, EUC encoding Japanese
EUC-KR	EUC_KR	KS C 5601, EUC encoding, Korean
GB18030	GB18030	Simplified Chinese, PRC standard
GB2312	EUC_CN	GB2312, EUC encoding, Simplified Chinese
GBK	GBK	GBK, Simplified Chinese
IBM-Thai	Cp838	IBM Thailand extended SBCS
IBM01140	Cp1140	Variant of Cp037 with Euro character
IBM01141	Cp1141	Variant of Cp273 with Euro character
IBM01142	Cp1142	Variant of Cp277 with Euro character
IBM01143	Cp1143	Variant of Cp278 with Euro character
IBM01144	Cp1144	Variant of Cp280 with Euro character
IBM01145	Cp1145	Variant of Cp284 with Euro character
IBM01146	Cp1146	Variant of Cp285 with Euro character
IBM01147	Cp1147	Variant of Cp297 with Euro character
IBM01148	Cp1148	Variant of Cp500 with Euro character
IBM01149	Cp1149	Variant of Cp871 with Euro character
IBM037	Cp037	USA, Canada (Bilingual, French), Netherlands, Portugal, Brazil, Australia
IBM1026	Cp1026	IBM Latin-5, Turkey
IBM1047	Cp1047	Latin-1 character set for EBCDIC hosts
IBM273	Cp273	IBM Austria, Germany

Canonical Name for <code>java.nio</code> API	Canonical Name for <code>java.io</code> API and <code>java.lang</code> API	Description
IBM277	Cp277	IBM Denmark, Norway
IBM278	Cp278	IBM Finland, Sweden
IBM280	Cp280	IBM Italy
IBM284	Cp284	IBM Catalan/Spain, Spanish Latin America
IBM285	Cp285	IBM United Kingdom, Ireland
IBM290	Cp290	EBCDIC-JP-kana, Japanese EBCDIC
IBM297	Cp297	IBM France
IBM300	Cp300	Japan DB EBCDIC, Japanese EBCDIC
IBM420	Cp420	IBM Arabic
IBM424	Cp424	IBM Hebrew
IBM500	Cp500	EBCDIC 500V1
IBM860	Cp860	MS-DOS Portuguese
IBM861	Cp861	MS-DOS Icelandic
IBM863	Cp863	MS-DOS Canadian French
IBM864	Cp864	PC Arabic
IBM865	Cp865	MS-DOS Nordic
IBM868	Cp868	MS-DOS Pakistan
IBM869	Cp869	IBM Modern Greek
IBM870	Cp870	IBM Multilingual Latin-2
IBM871	Cp871	IBM Iceland
IBM918	Cp918	IBM Pakistan (Urdu)
ISO-2022-CN	ISO2022CN	GB2312 and CNS11643 in ISO 2022 CN form, Simplified and Traditional Chinese (conversion to Unicode only)
ISO-2022-JP	ISO2022JP	JIS X 0201, 0208, in ISO 2022 form, Japanese
ISO-2022-KR	ISO2022KR	ISO 2022 KR, Korean
ISO-8859-3	ISO8859_3	Latin Alphabet No. 3
ISO-8859-6	ISO8859_6	Latin/Arabic Alphabet
ISO-8859-8	ISO8859_8	Latin/Hebrew Alphabet
JIS_X0201	JIS_X0201	JIS X 0201
JIS_X0212-1990	JIS_X0212-1990	JIS X 0212
Shift_JIS	SJIS	Shift-JIS, Japanese
TIS-620	TIS620	TIS620, Thai
windows-1255	Cp1255	Windows Hebrew
windows-1256	Cp1256	Windows Arabic
windows-1258	Cp1258	Windows Vietnamese

Canonical Name for <code>java.nio</code> API	Canonical Name for <code>java.io</code> API and <code>java.lang</code> API	Description
<code>windows-31j</code>	<code>MS932</code>	Windows Japanese
<code>x-Big5-Solaris</code>	<code>Big5_Solaris</code>	Big5 with seven additional Hanzi ideograph character mappings for the Solaris <code>zh_TW.BIG5</code> locale
<code>x-euc-jp-linux</code>	<code>EUC_JP_LINUX</code>	JISX 0201, 0208, EUC encoding Japanese
<code>x-EUC-TW</code>	<code>EUC_TW</code>	CNS11643 (Plane 1-7,15), EUC encoding, Traditional Chinese
<code>x-eucJP-Open</code>	<code>EUC_JP_Solaris</code>	JISX 0201, 0208, 0212, EUC encoding Japanese
<code>x-IBM1006</code>	<code>Cp1006</code>	IBM AIX Pakistan (Urdu)
<code>x-IBM1025</code>	<code>Cp1025</code>	IBM Multilingual Cyrillic: Bulgaria, Bosnia, Herzegovina, Macedonia (FYR)
<code>x-IBM1046</code>	<code>Cp1046</code>	IBM Arabic - Windows
<code>x-IBM1097</code>	<code>Cp1097</code>	IBM Iran (Farsi)/Persian
<code>x-IBM1098</code>	<code>Cp1098</code>	IBM Iran (Farsi)/Persian (PC)
<code>x-IBM1112</code>	<code>Cp1112</code>	IBM Latvia, Lithuania
<code>x-IBM1122</code>	<code>Cp1122</code>	IBM Estonia
<code>x-IBM1123</code>	<code>Cp1123</code>	IBM Ukraine
<code>x-IBM1124</code>	<code>Cp1124</code>	IBM AIX Ukraine
<code>x-IBM1381</code>	<code>Cp1381</code>	IBM OS/2, DOS People's Republic of China (PRC)
<code>x-IBM1383</code>	<code>Cp1383</code>	IBM AIX People's Republic of China (PRC)
<code>x-IBM33722</code>	<code>Cp33722</code>	IBM-eucJP - Japanese (superset of 5050)
<code>x-IBM834</code>	<code>Cp834</code>	IBM EBCDIC DBCS-only Korean
<code>x-IBM856</code>	<code>Cp856</code>	IBM Hebrew
<code>x-IBM875</code>	<code>Cp875</code>	IBM Greek
<code>x-IBM921</code>	<code>Cp921</code>	IBM Latvia, Lithuania (AIX, DOS)
<code>x-IBM922</code>	<code>Cp922</code>	IBM Estonia (AIX, DOS)
<code>x-IBM930</code>	<code>Cp930</code>	Japanese Katakana-Kanji mixed with 4370 UDC, superset of 5026
<code>x-IBM933</code>	<code>Cp933</code>	Korean Mixed with 1880 UDC, superset of 5029
<code>x-IBM935</code>	<code>Cp935</code>	Simplified Chinese Host mixed with 1880 UDC, superset of 5031

Canonical Name for <code>java.nio</code> API	Canonical Name for <code>java.io</code> API and <code>java.lang</code> API	Description
x-IBM937	Cp937	Traditional Chinese Host mixed with 6204 UDC, superset of 5033
x-IBM939	Cp939	Japanese Latin Kanji mixed with 4370 UDC, superset of 5035
x-IBM942	Cp942	IBM OS/2 Japanese, superset of Cp932
x-IBM942C	Cp942C	Variant of Cp942
x-IBM943	Cp943	IBM OS/2 Japanese, superset of Cp932 and Shift-JIS
x-IBM943C	Cp943C	Variant of Cp943
x-IBM948	Cp948	OS/2 Chinese (Taiwan) superset of 938
x-IBM949	Cp949	PC Korean
x-IBM949C	Cp949C	Variant of Cp949
x-IBM950	Cp950	PC Chinese (Hong Kong, Taiwan)
x-IBM964	Cp964	AIX Chinese (Taiwan)
x-IBM970	Cp970	AIX Korean
x-ISCII91	ISCII91	ISCII91 encoding of Indic scripts
x-ISO2022-CN-CNS	ISO2022_CN_CNS	CNS11643 in ISO 2022 CN form, Traditional Chinese (conversion from Unicode only)
x-ISO2022-CN-GB	ISO2022_CN_GB	GB2312 in ISO 2022 CN form, Simplified Chinese (conversion from Unicode only)
x-iso-8859-11	x-iso-8859-11	Latin/Thai Alphabet
x-JIS0208	x-JIS0208	JIS X 0208
x-JISAutoDetect	JISAutoDetect	Detects and converts from Shift-JIS, EUC-JP, ISO 2022 JP (conversion to Unicode only)
x-Johab	x-Johab	Korean, Johab character set
x-MacArabic	MacArabic	Macintosh Arabic
x-MacCentralEurope	MacCentralEurope	Macintosh Latin-2
x-MacCroatian	MacCroatian	Macintosh Croatian
x-MacCyrillic	MacCyrillic	Macintosh Cyrillic
x-MacDingbat	MacDingbat	Macintosh Dingbat
x-MacGreek	MacGreek	Macintosh Greek
x-MacHebrew	MacHebrew	Macintosh Hebrew

Canonical Name for <code>java.nio</code> API	Canonical Name for <code>java.io</code> API and <code>java.lang</code> API	Description
<code>x-MacIceland</code>	<code>MacIceland</code>	Macintosh Iceland
<code>x-MacRoman</code>	<code>MacRoman</code>	Macintosh Roman
<code>x-MacRomania</code>	<code>MacRomania</code>	Macintosh Romania
<code>x-MacSymbol</code>	<code>MacSymbol</code>	Macintosh Symbol
<code>x-MacThai</code>	<code>MacThai</code>	Macintosh Thai
<code>x-MacTurkish</code>	<code>MacTurkish</code>	Macintosh Turkish
<code>x-MacUkraine</code>	<code>MacUkraine</code>	Macintosh Ukraine
<code>x-MS950-HKSCS</code>	<code>MS950_HKSCS</code>	Windows Traditional Chinese with Hong Kong extensions
<code>x-mswin-936</code>	<code>MS936</code>	Windows Simplified Chinese
<code>x-PCK</code>	<code>PCK</code>	Solaris version of <code>Shift_JIS</code>
<code>x-SJIS_0213</code>	<code>x-SJIS_0213</code>	<code>Shift_JISX0213</code>
<code>x-windows-50220</code>	<code>Cp50220</code>	Windows Codepage 50220 (7- bit implementation)
<code>x-windows-50221</code>	<code>Cp50221</code>	Windows Codepage 50221 (7- bit implementation)
<code>x-windows-874</code>	<code>MS874</code>	Windows Thai
<code>x-windows-949</code>	<code>MS949</code>	Windows Korean
<code>x-windows-950</code>	<code>MS950</code>	Windows Traditional Chinese
<code>x-windows-iso2022jp</code>	<code>x-windows-iso2022jp</code>	Variant ISO-2022-JP (MS932 based)

4

Supported Calendars

The core of the Date-Time API is the [java.time](#) package. The classes defined in `java.time` base their calendar system on the ISO calendar, which is the world standard for representing date and time. The ISO calendar follows the proleptic Gregorian rules. There are also non-ISO calendars predefined in `java.time.chrono` package: the Japanese, Hijrah, Minguo, and Thai Buddhist calendars. For more about the Date-Time API, see the [Internationalization Trail](#) in the Java Tutorials.

5

Supported Fonts

The set of supported fonts varies between different implementations of the Java platform. For the terminology used, see the [Font](#) class description.

- [Support for Physical Fonts](#)
- [Support for Logical Fonts](#)
- [The Lucida Fonts](#)

Support for Physical Fonts

The JRE supports TrueType and PostScript Type 1 fonts.

Physical fonts need to be installed in locations known to the Java runtime environment. The JRE looks in two locations: the `lib/fonts` directory within the JRE itself, and the normal font location(s) defined by the host operating system. If fonts with the same name exist in both locations, the one in the `lib/fonts` directory is used.

You can add physical fonts that use a supported font technology by installing them in a way supported by the host operating system. The recommended location to add per-user fonts on Solaris or Linux is the `$HOME/.fonts` directory which is searched by the platform's `libfontconfig`, and which is in turn used by the JRE.

Do not make any modifications under the `lib` directory within the JRE, as this is not supported since JDK 9.

Support for Logical Fonts

Logical font names are mapped to physical fonts in implementation dependent ways. Typically one logical font name maps to several physical fonts in order to cover a large range of characters. The JRE uses font configuration files to define the mapping; see [Font Configuration Files](#).

The Lucida Fonts

The Oracle JRE includes several physical fonts of the "Lucida" design family. These fonts are physical fonts, but since they come with the JRE, they don't depend on the host operating system. These fonts are also licensed for use in other implementations of the Java Platform. Using them provides the benefit of a consistent look and feel across platforms and implementations for a large set of languages.

There are three different type families: "Lucida Sans", "Lucida Sans Typewriter", and "Lucida Bright". Each family has plain, bold, italic, and bold-italic styles. Not all of these are present in all JRE implementations. For example, the default download bundle of the JRE for Windows only contains the Lucida Sans plain font (but note that application developers can include all Lucida fonts with a JRE that they redistribute with an application).

The following table shows which Unicode character blocks are covered by each font family:

Unicode Block	Lucida Sans	Lucida Sans Typewriter	Lucida Bright
Basic Latin	yes	yes	yes
Latin-1 Supplement	yes	yes	yes
Latin Extended-A	yes	yes	yes
Latin Extended-B	partial	partial	partial
IPA Extensions	partial	no	no
Spacing Modifier Letter	partial	partial	partial
Combining Diacritical Marks	partial	no	partial
Greek	yes	yes	yes
Cyrillic	partial	partial	partial
Hebrew	yes	yes	no
Arabic	partial	partial	partial
Devanagari	yes	no	no
Thai	yes	yes	no
General Punctuation	partial	partial	partial
Superscripts and Subscripts	partial	partial	partial
Currency Symbols	partial	partial	partial
Letterlike Symbols	partial	partial	partial
Number Forms	partial	partial	no
Arrows	partial	partial	partial
Mathematical Operators	yes	partial	yes
Enclosed Alphanumerics	partial	no	no
Box Drawings	partial	partial	no
Block Elements	partial	partial	no
Geometric Shapes	partial	partial	partial
Miscellaneous Symbols	partial	partial	no
Dingbats	yes	no	no
Alphabetic Presentation Forms	partial	partial	no
Arabic Presentation Forms-A	partial	partial	partial
Arabic Presentation Forms-B	yes	yes	yes

Note that of the writing systems that are generally fully supported by the JRE, the Lucida fonts do not support Chinese (Simplified), Chinese (Traditional), Japanese, and

Korean. See the "Supported Locales" page for this release on [Oracle Technology Network](#).

The fonts are installed in the Java SE Runtime Environment's `lib/fonts` directory as the following files (not all of them may be present):

```
LucidaSansDemiBold.ttf  
LucidaSansRegular.ttf  
LucidaTypewriterBold.ttf  
LucidaTypewriterRegular.ttf  
LucidaBrightDemiBold.ttf  
LucidaBrightDemiItalic.ttf  
LucidaBrightItalic.ttf  
LucidaBrightRegular.ttf
```

6

Font Configuration Files

The Java Platform defines five logical font names that every implementation must support: `Serif`, `SansSerif`, `Monospaced`, `Dialog`, and `DialogInput`. These logical font names are mapped to physical fonts in implementation dependent ways.

One way the Oracle JDK maps logical font names to physical fonts is by using *font configuration* files. There may be several files to support different mappings depending on the host operating system version. The files are distributed with the JDK installation. You can edit or create your own font configuration files to adjust the mappings to your particular system setup, however these must be placed in `conf/` `fonts`, and are subject to implementation notes discussed below.

Font configuration files come in two formats: a properties format and a binary format. The properties format is described in detail in this document and can be used for user-defined configurations. The binary format is undocumented and used only for the JDK's predefined configurations; the corresponding files in properties format are available for reference as files with the `.properties.src` extension.

Supported Platforms

Font configuration files are implementation dependent. Not all implementations of the Java Platform use them, and the format and content vary between different runtime environments as well as between releases. The macOS implementation does not use font configuration files, as the mapping is hard coded in the source and cannot be changed in any way.

The Oracle JDK supports font configuration files on the host operating system as follows:

- For Windows, font configuration files are required.
- For macOS, font configuration files are unsupported.
- For Linux and Solaris: the Oracle JDK is moving away from providing custom font configuration files on Linux platforms, as they are difficult to keep up to date across distributions and versions. A distribution that has control over the fonts on the system can continue to provide this custom file. If the JRE finds a custom file that exactly matches the distribution and version it will use it. If no exact match is found, the JRE dynamically creates the file at runtime. These generated files are placed in a location determined by the implementation. They should be considered implementation internal: they are not user editable and do not follow the syntax as described in this specification.

Loading Font Configuration Files

The JDK places any files that it provides in `$JDKHOME/lib`. Do not modify that location. Instead, put any updates or custom versions of the configuration files in `$JDKHOME/conf/fonts`.

On platforms that support font configuration files, the runtime will look first in `$JDKHOME/conf/fonts`. In other words, a user-supplied file is preferred if it is a match.

The font configuration file for a host operating system is located as follows:

- *JavaHome* - the JDK directory, as given by the `java.home` system property.
- *OS* - a string identifying an operating system variant:
 - For Windows, empty.
 - For Solaris, empty.
 - For Linux, "RedHat", "SuSE", etc.
- *Version* - a string identifying the operating system version.

The runtime uses the first of the following files it finds:

```
JavaHome/lib/fontconfig.OS.Version.properties
JavaHome/lib/fontconfig.OS.Version.bfc
JavaHome/lib/fontconfig.OS.properties
JavaHome/lib/fontconfig.OS.bfc
JavaHome/lib/fontconfig.Version.properties
JavaHome/lib/fontconfig.Version.bfc
JavaHome/lib/fontconfig.properties
JavaHome/lib/fontconfig.bfc
```

Files with a `.properties` suffix are assumed to be properties files as specified by the [Properties](#) class and are loaded through that class. Files without this suffix are assumed to be in binary format.

Names Used in Font Configuration Files

Throughout the font configuration files, a number of different names are used:

- *LogicalFontName* - one of the five logical font names: `serif`, `sansserif`, `monospaced`, `dialog`, and `dialoginput`. In font configuration files, these names are always in lowercase.
- *StyleName* - one of the four standard font styles: `plain`, `bold`, `italic`, and `bolditalic`. Again, these names are always in lowercase.
- *PlatformFontName* - the name of a physical font, in a format typically used on the platform:
 - On Windows, a font face name, such as "Courier New" or "\uad74\uub9bc".
 - On Solaris and Linux, an xld name, such as "-monotype-times new roman-regular-r---*-%d-*-*p-*-*iso8859-1". Note that "%d" is used for the font size - the actual font size is filled in at runtime.
- *CharacterSubsetName* - a name for a subset of the Unicode character set which certain component fonts can render. For Windows, the following names are predefined: `alphabetic`, `arabic`, `chinese-ms936`, `chinese-gb18030`, `chinese-ms950`, `chinese-hkscs`, `cyrillic-iso8859-5`, `cyrillic-cp1251`, `cyrillic-koi8-r`, `devanagari`, `dingbats`, `greek`, `hebrew`, `japanese`, `korean`, `latin`, `symbol`, `thai`. For Solaris and Linux, the following names are predefined: `arabic`, `chinese-gb2312`, `chinese-gbk`, `chinese-gb18030-0`, `chinese-gb18030-1`, `chinese-cns11643-1`, `chinese-cns11643-2`, `chinese-cns11643-3`, `chinese-big5`, `chinese-hkscs`, `cyrillic`, `devanagari`, `dingbats`, `greek`, `hebrew`, `japanese-x0201`, `japanese-x0208`, `japanese-`

x0212, korean, korean-johab, latin-1, latin-2, latin-4, latin-5, latin-7, latin-9, symbol, thai. A font configuration file may define additional names to identify additional character subsets.

- *Encoding* - the canonical name of the default encoding, as provided by `java.nio.charset.Charset.defaultCharset().name()`.
- *Language* - the language of the initial default locale.
- *Country* - the country of the initial default locale.

Properties for All Platforms

Properties that are applicable to all platforms enable you to specify the font configuration format version, component font mappings, search sequences, exclusion ranges, proportional fonts, font file names, and appended font path.

Version Property

The version property identifies the font configuration format version. This document specifies version 1.

The complete property has the form:

```
version=1
```

Component Font Mappings

Component font mapping properties describe which physical font to use to render characters from a given character subset with a given logical font in a given style.

The keys have the forms:

```
allfonts.CharacterSubsetName  
LogicalFontName.StyleName.CharacterSubsetName
```

The first form is used if the same font is used for a character subset independent of logical font and style (in this case, the font rendering engines apply algorithmic styles to the font). The second form is used if different physical fonts are used for a character subset for different logical fonts and styles. In this case, properties must be specified for each combination of logical font and style, so 20 properties for one character subset. If a property of the first form is present for a character subset, then properties of the second form for the same character subset are ignored.

The values are platform font names, as described in [Names Used in Font Configuration Files](#).

Since the character subsets supported by given fonts often overlap, separate search sequence properties are used to define in which order to try the fonts when rendering a character.

Search Sequences

The Java runtime uses sequence properties to determine search sequences for the five logical fonts. However, a font configuration file may specify properties that are specific to a combination of encoding, language, and country, and the runtime will then use a lookup to determine the search sequence property for each logical font.

The keys have the form:

```
sequence.allfonts.Encoding.Language.Country  
sequence.LogicalFontName.Encoding.Language.Country  
sequence.allfonts.Encoding.Language  
sequence.LogicalFontName.Encoding.Language  
sequence.allfonts.Encoding  
sequence.LogicalFontName.Encoding  
sequence.allfonts  
sequence.LogicalFontName
```

The `allfonts` forms are used if the sequence is used for all five logical fonts. The forms specifying logical font names are used if different sequences are used for different logical fonts.

For each logical font, the Java runtime uses the property value with the first of the above keys. This property determines the primary search sequence for the logical font.

The file may also define a single fallback search sequence. The key for the fallback search sequence property is:

```
sequence.fallback
```

The values of all search sequence properties have the form:

```
SearchSequenceValue:  
  CharacterSubsetName  
  CharacterSubsetName , SearchSequenceValue
```

The primary search sequence properties specify character subset names for required fonts, which are used for both AWT and 2D font rendering. The fallback search sequence property gives character subset names for optional fonts, which are used as fallbacks for all logical fonts, but only for 2D font rendering. The runtime automatically adds the Lucida Sans Regular font as a fallback font for 2D rendering if it's not already specified. If the runtime environment has a `lib/fonts/fallback` directory that contains valid TrueType or Type 1 fonts, the runtime automatically adds these fonts as fallback fonts for 2D rendering. On Windows, if there is a system EUDC (End User Defined Characters) font registered with Windows, the runtime automatically adds this font as well as a fallback font for 2D rendering.

The sequence properties determine in which sequence component fonts are tried to render a given character. For example, given the following properties:

```
sequence.monospaced=japanese,alphabetic  
sequence.fallback=korean  
monospaced.plain.alphabetic=Arial  
monospaced.plain.japanese=MSGothic  
monospaced.plain.korean=Gulim
```

The runtime will first attempt to render a character with the MSGothic font. If that font doesn't provide a glyph for the character, it will attempt the Arial font. For 2D rendering, it will also try the Gulim and the Lucida Sans Regular font as well as any TrueType or Type 1 fonts in the runtime's `lib/fonts/fallback` directory. For 2D rendering on Windows, if there is a system EUDC font registered with Windows, the runtime will also try this EUDC font.

When calculating font metrics for a logical font without reference to a string, only the required fonts are taken into consideration. For the example above, the [FontMetrics.getMaxDescent](#) method would return results based on the MSGothic and Arial fonts, but not the Gulim and Lucida Sans fonts. In this way, simple user

interface elements such as buttons, which sometimes calculate their size based on font metrics, are not affected by an extended list of component fonts which their labels usually don't use. On the other hand, text components typically calculate metrics based on the text they contain and thus will obtain correct results.

The sequence properties that the runtime obtains for the five logical fonts should list the same character subsets, but may list them in different order.

Exclusion Ranges

The exclusion range properties specify Unicode character ranges which should be excluded from being rendered with the fonts corresponding to a given character subset. This is used if a font with a large character repertoire needs to be placed early in the search sequence (for example, for performance reasons), but some characters that it supports should be drawn with a different font instead. These properties are optional, so there's at most one per character subset.

The keys have the form:

```
exclusion.CharacterSubsetName
```

The values have the form:

```
ExclusionRangeValue:  
    Range  
    Range , ExclusionRangeValue
```

```
Range:  
    Char - Char
```

```
Char:  
    HexDigit HexDigit HexDigit HexDigit  
    HexDigit HexDigit HexDigit HexDigit HexDigit HexDigit
```

A *Char* is a Unicode character represented as a hexadecimal value.

Proportional Fonts

The proportional font properties describe the relationship between proportional and non-proportional variants of otherwise equivalent fonts. These properties are used to implement preferences specified by the [GraphicsEnvironment.preferProportionalFonts](#) method.

The keys have the form:

```
proportional.PlatformFontName
```

Space characters in the platform font name must be replaced with underscore characters (_).

The values have the form:

```
PlatformFontName
```

In values, space characters are left unmodified.

Each property indicates that the font named in the value is the proportional equivalent of the font named in the key, and also that the font named in the key is the non-proportional equivalent of the font named in the value.

Font File Names

Font file name properties provide the names of the files containing the physical fonts used in the font configuration file. File names are required for all physical fonts on Windows and recommended for all physical fonts on Solaris and Linux.

The keys have the form:

filename.PlatformFontName

Space characters in the platform font name must be replaced with underscore characters (_).

The values are the file names of the files containing the fonts. On Windows, simple file names are used; and the runtime environment looks for each file first in its own `lib/fonts` directory, then in the Windows fonts directory. On Solaris and Linux, absolute path names, path names starting with "`$JRE_LIB_FONTS`" for the runtime environment's own `lib/fonts` directory, or `xlfd` names are used.

Appended Font Path

The Java runtime can automatically determine a number of directories that contain font files, such as its own `lib/fonts` directory or the Windows fonts folder. Additional directories can be specified to be appended to the font path.

The key has the form:

appendedfontpath

The value has the form:

AppendedFontPathValue:
Directory
Directory PathSeparator AppendedFontPathValue

The path separator is the platform dependent value of `java.io.File.pathSeparator`.

Properties for Windows

There are no platform-specific properties for Windows. However, there is a special form of the character subset name used in search sequences. The name "alphabetic" can take a suffix indicating the character encoding associated with the subset:

alphabetic
alphabetic/default
alphabetic/1252

This information is only used for AWT, not for 2D. The `/default` suffix restricts use of the component fonts for this character subset to the character set of the default encoding; the `/1252` suffix to the Windows-1252 character set. For accessing component font mappings and exclusion ranges, the character encoding suffix is omitted. For all other character subsets, the AWT character encoding is determined internally by the Java runtime.

Property for Solaris and Linux

The only property that is specific to Solaris and Linux is the AWT font path, which identifies platform directories that should be added to the X11 server font path.

The keys have the form:

`awtfontpath.CharacterSubsetName`

The values have the form:

```
AWTFontPathValue:  
  Directory  
  Directory : AWTFontPathValue
```

The directories must be valid X11 font directories. The Java runtime ensures that the directories for all character subsets of a primary search sequence found by the search sequence lookup (see [Search Sequences](#)) are part of the X11 font path. The implementation assumes that all logical fonts use the same set of character subsets for a given environment of encoding, language, and country.